



# ECE 6775

## High-Level Digital Design Automation

### Fall 2024

## Course Overview

Zhiru Zhang

School of Electrical and Computer Engineering



Cornell University



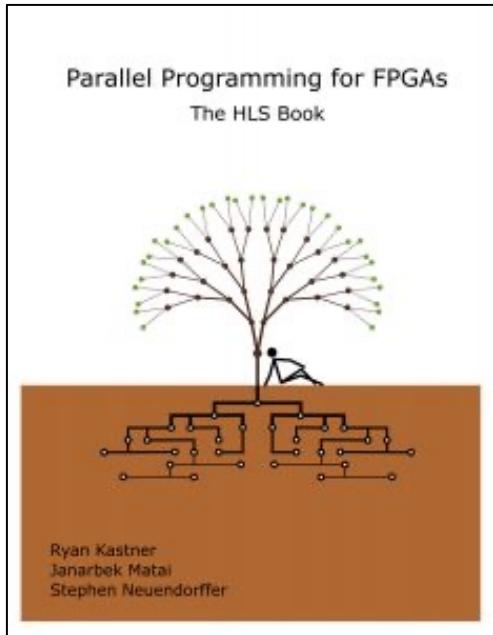
# Agenda

- ▶ Important logistics
- ▶ Course motivation
- ▶ More course organization

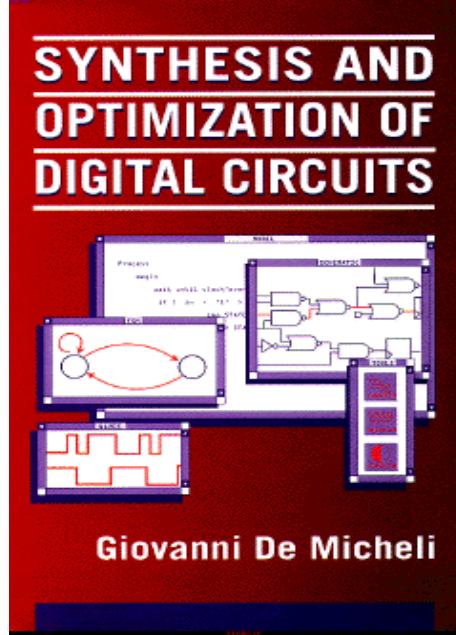
# Class Resources

- ▶ Course website
  - <https://www.csl.cornell.edu/courses/ece6775>
  - Lectures slides, handouts, and other readings
- ▶ Ed Discussion
  - Announcements and Q&A
- ▶ CMSX: course management system
  - Assignments and grades
  - Electronic submissions required

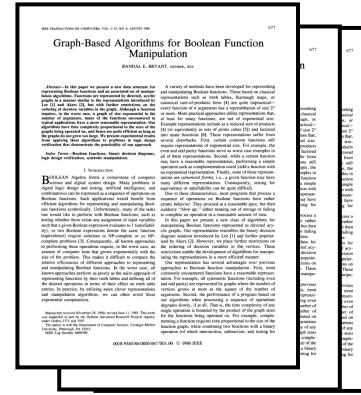
# Course Texts



**e-book available online**



Get 1st edition  
Overhead slides  
available online



## Selected papers & software manuals

# Seeking Help After Class

- ▶ Ed Discussion
  - Questions on lectures, assignments, projects, etc.
  - Monitored by course staff
- ▶ Instructor office hours (online)
  - Thursday 4:30-5:30pm, Zoom link posted on Ed
- ▶ Email instructor for personal issues/appointment
- ▶ PhD TAs:
  - Yixiao Du (yd383), Head TA
  - Andrew Butt (atb78), part-time

# Grading Scheme

- ▶ Class participation (4%)
  - Asking & answering questions during lectures
  - Contributing to discussions on Ed
- ▶ Paper readings (5%)
- ▶ Quizzes (6%)
- ▶ Midterm exam (20%)
- ▶ Assignments (30%)
- ▶ Final project (35%)

# What this Course is About

## *Hardware/Software Co-Design*

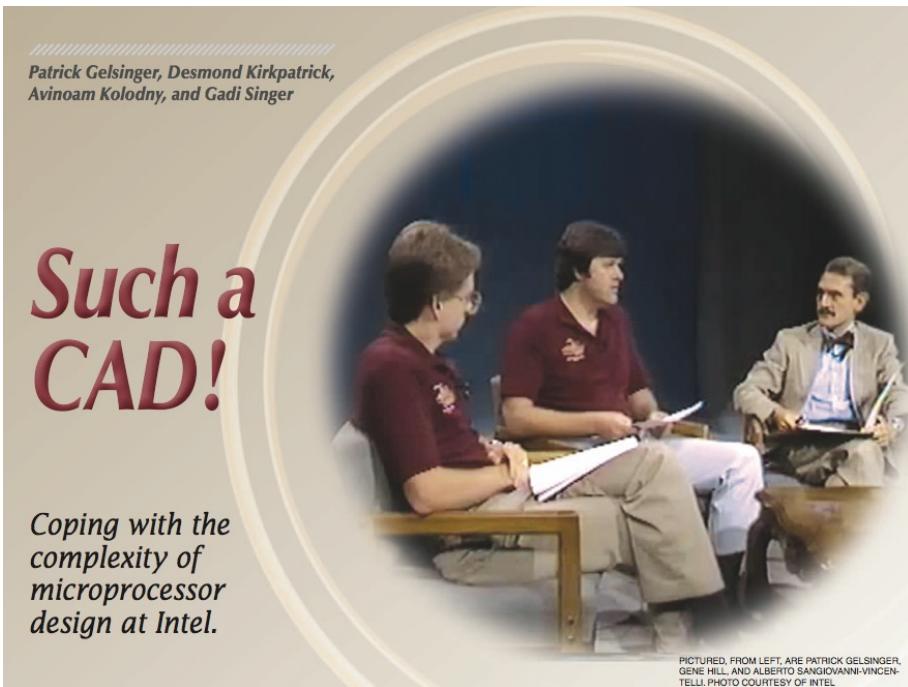
- ▶ Specify applications/algorithms in software programs
- ▶ Synthesize software descriptions into special-purpose hardware components, namely, *accelerators*
  - Perform manual source-level code optimizations
  - Utilize automatic compilation & synthesis optimizations
  - Explore performance-cost trade-offs
- ▶ Realize the synthesized accelerators on FPGAs

# This Course Dives into EDA

## *Electronic Design Automation*

- ▶ A general methodology for refining **a high-level description down to a detailed physical implementation** for designs ranging from
  - integrated circuits (including system-on-chips),
  - printed circuit boards (PCBs),
  - and electronic systems
- ▶ **Modeling, synthesis, and verification** at every level of abstraction

# Significance of EDA



Coping with the complexity of microprocessor design at Intel.

Patrick Gelsinger, Desmond Kirkpatrick, Avinoam Kolodny, and Gadi Singer

**Such a CAD!**

PICTURED, FROM LEFT, ARE PATRICK GELSINGER, GENE HILL, AND ALBERTO SANGIOVANNI-VINCENTELLI. PHOTO COURTESY OF INTEL

Patrick Gelsinger, Desmond Kirkpatrick, Avinoam Kolodny, and Gadi Singer. "Such a CAD!" *IEEE Solid-State Circuits Magazine*, 2010.

TABLE 1. INTEL PROCESSORS, 1971–1993.

PROCESSOR	INTRO DATE	PROCESS	TRANSISTORS	FREQUENCY
4004	1971	10 $\mu$ m	2,300	108 KHz
8080	1974	6 $\mu$ m	6,000	2 MHz
8086	1978	3 $\mu$ m	29,000	10 MHz
80286	1982	1.5 $\mu$ m	134,000	12 MHz
80386	1985	1.5 $\mu$ m	275,000	16 MHz
Intel 486 DX	1989	1 $\mu$ m	1.2 M	33 MHz
Pentium	1993	0.8 $\mu$ m	3.1 M	60 MHz

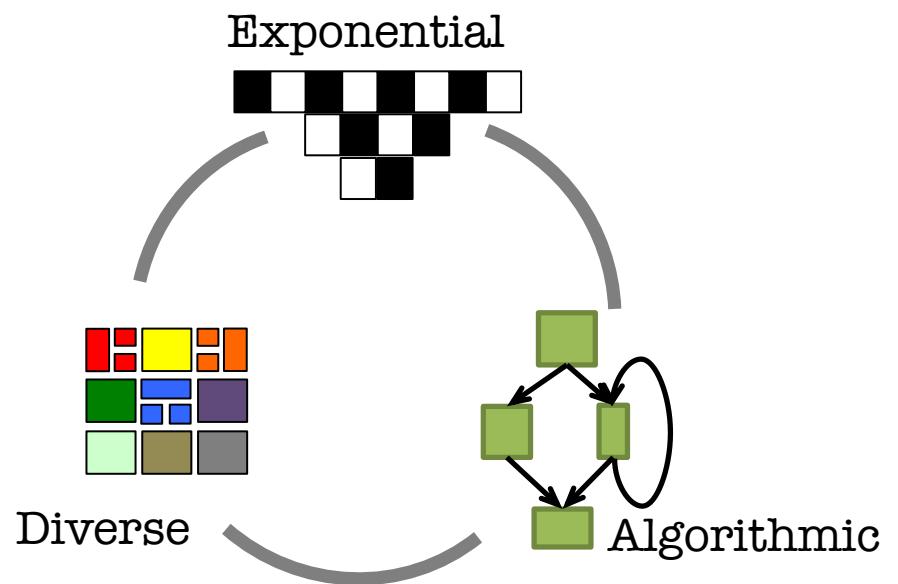
“ This incredible growth rate could not be achieved by hiring an exponentially growing number of design engineers. It was fulfilled by adopting new design methodologies and by introducing innovative design automation software at every processor generation. ”

# E-D-A: My Other Interpretation

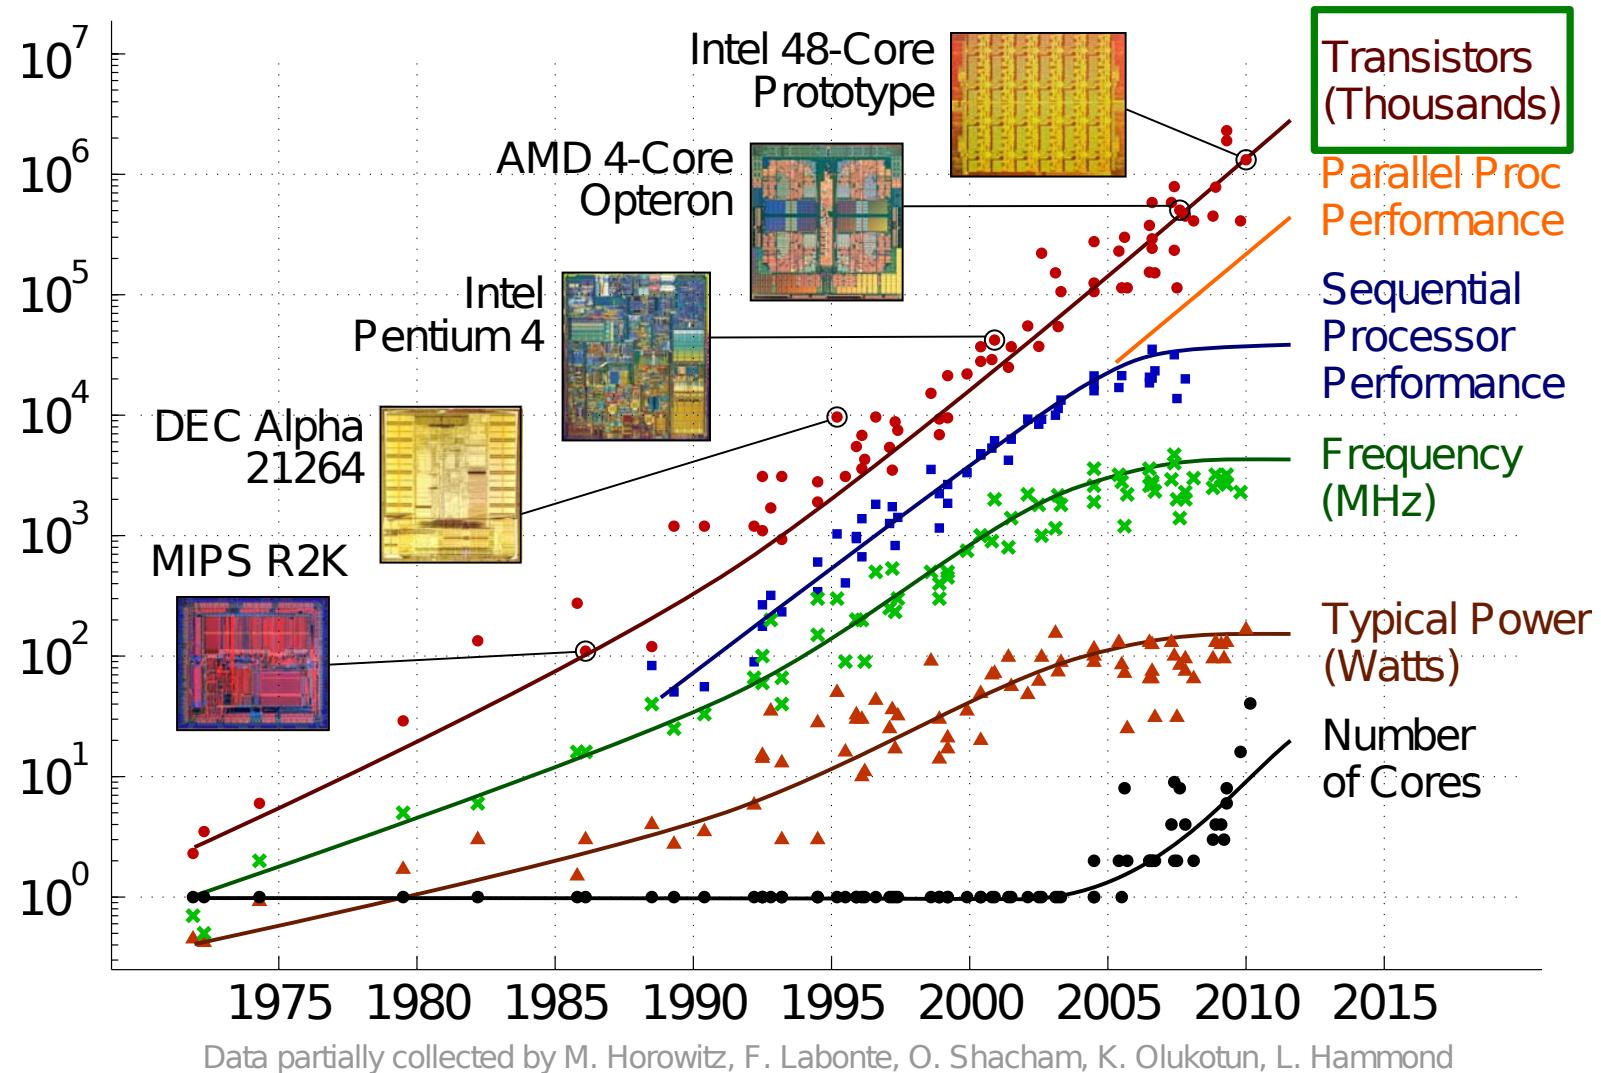
**Exponential**  
in complexity (or **Extreme scale**)

**Diverse**  
increasing system heterogeneity  
multi-disciplinary

**Algorithmic**  
intrinsically computational

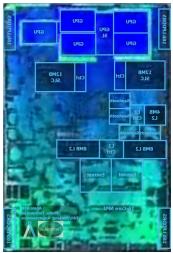


# Exponential: Moore's Law

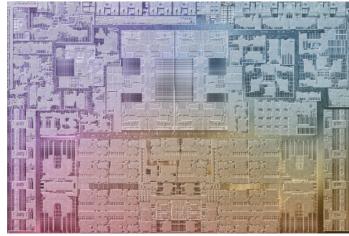


[Figure credit: Christopher Batten, Cornell]

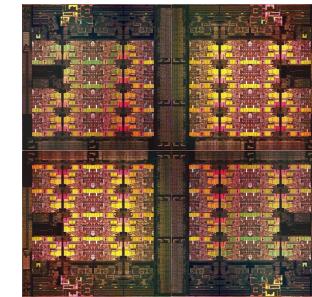
# Era of Billion-Transistor Chips



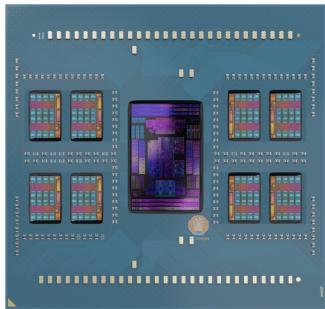
Apple A16  
~16B transistors



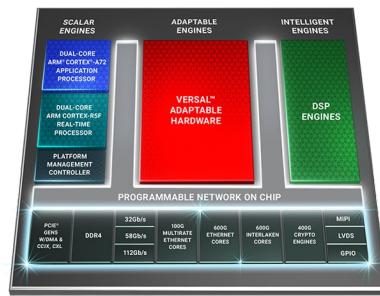
Apple M2 Pro  
~40B transistors



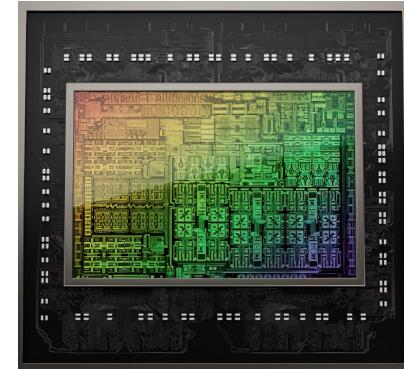
Intel Sapphire Rapids  
(quad-chip module)  
~48B transistors



AMD EPYC Bergamo  
(9-chip module)  
~82B transistors

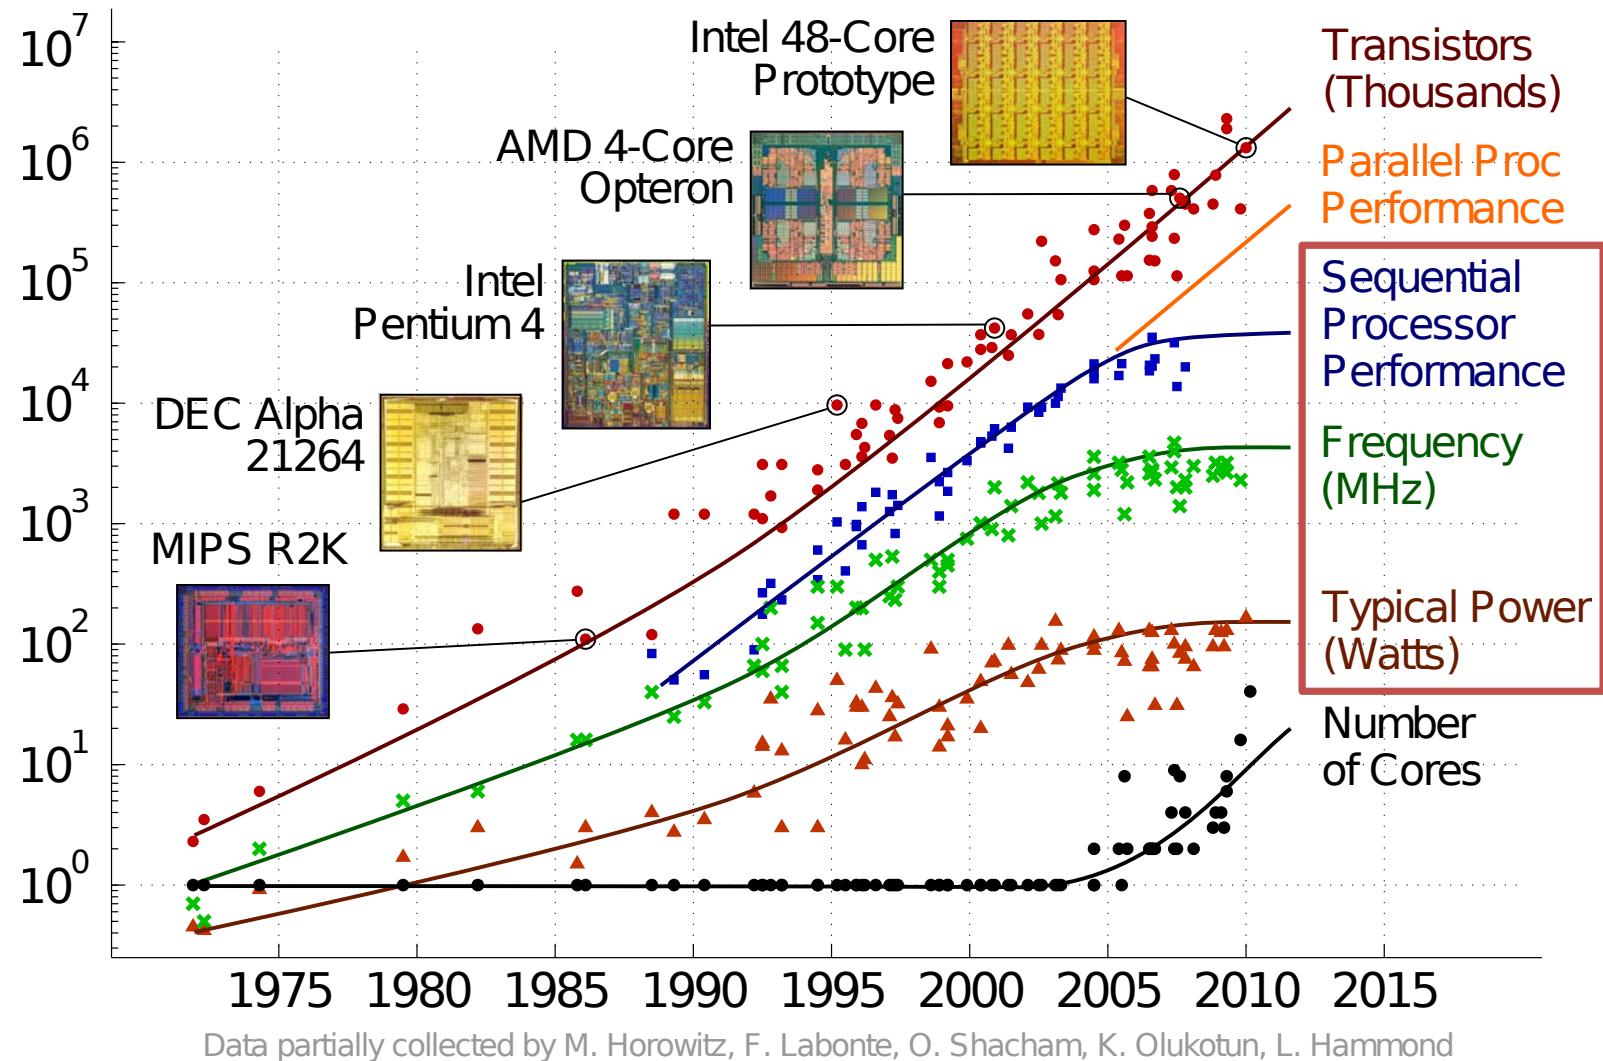


AMD (Xilinx) Versal Premium  
~92B transistors



NVIDIA Blackwell B200  
Grace Hopper Superchip  
~208B transistors

# End of Dennard Scaling: Power Becomes the Limiting Factor



[Figure credit: Christopher Batten, Cornell]

# Power-Constrained Modern Computers



$$Power = \frac{Energy}{Second} = \frac{Energy}{Op} \times \frac{Ops}{Second}$$

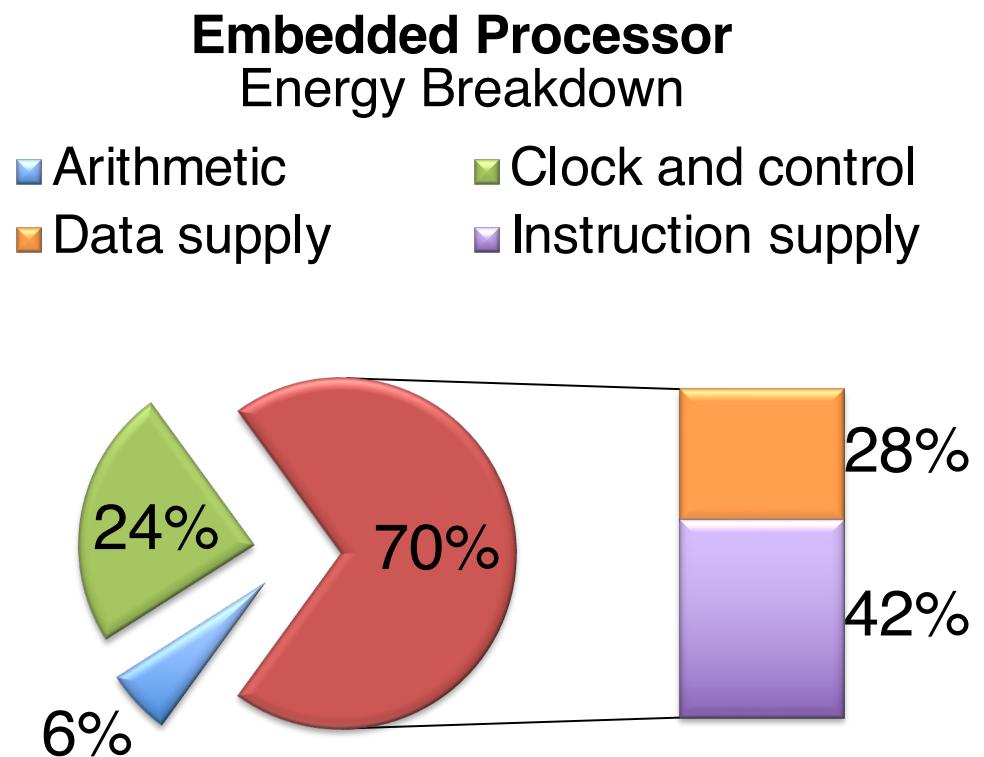
A red arrow points down to the term  $\frac{Energy}{Op}$ . A green arrow points up to the term  $\frac{Ops}{Second}$ .

To increase performance (Ops/sec) in a power-constrained regime, energy per operation must decrease—in other words, **energy efficiency** (Ops/Joule) **needs to improve!**

# Inefficiency of General-Purpose Computing

- ▶ Typical energy overhead (or “tax”) for every 10pJ arithmetic operations
  - 70pJ on instruction supply
  - 47pJ on data supply

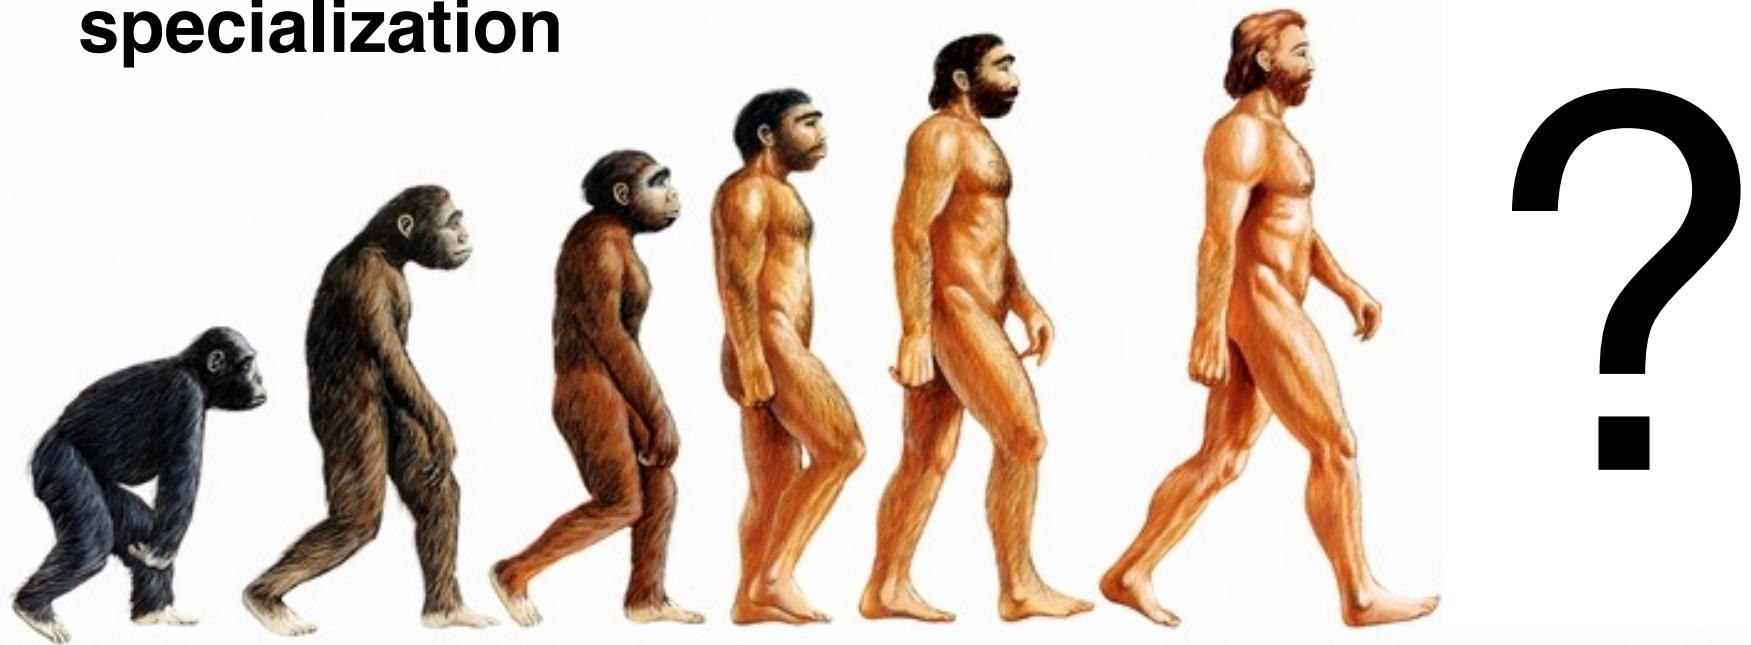
Also, only 59% of the instructions are arithmetic



[source: Dally et al. Efficient Embedded Computing, IEEE'08]

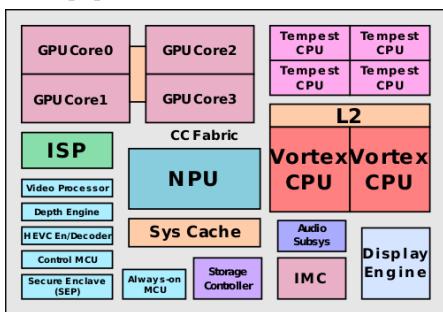
# Advance of Civilization

- ▶ For humans, Moore's Law scaling of the brain has ended a long time ago
  - Number of neurons and their firing rate did not change significantly
- ▶ Remarkable advancement of civilization via **specialization**

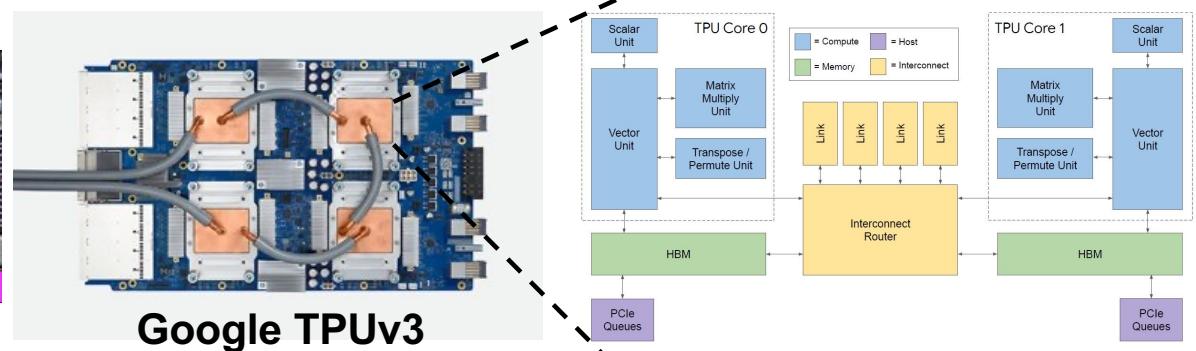
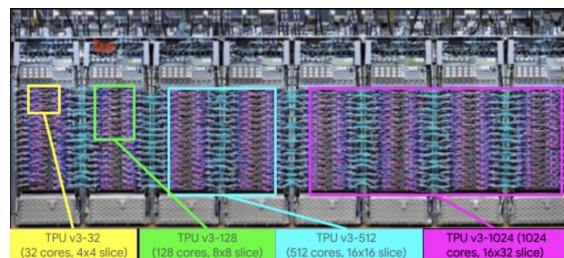
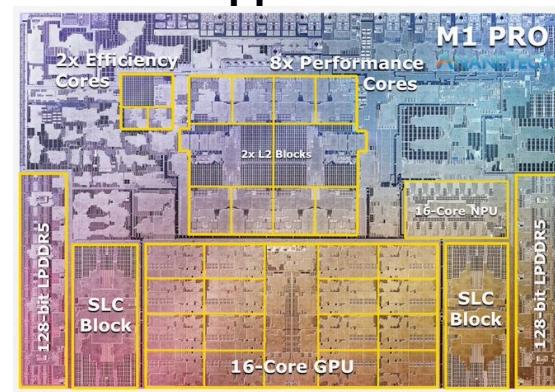


# Diverse: Era of Hardware Heterogeneity

Apple 12 (iPhone X)



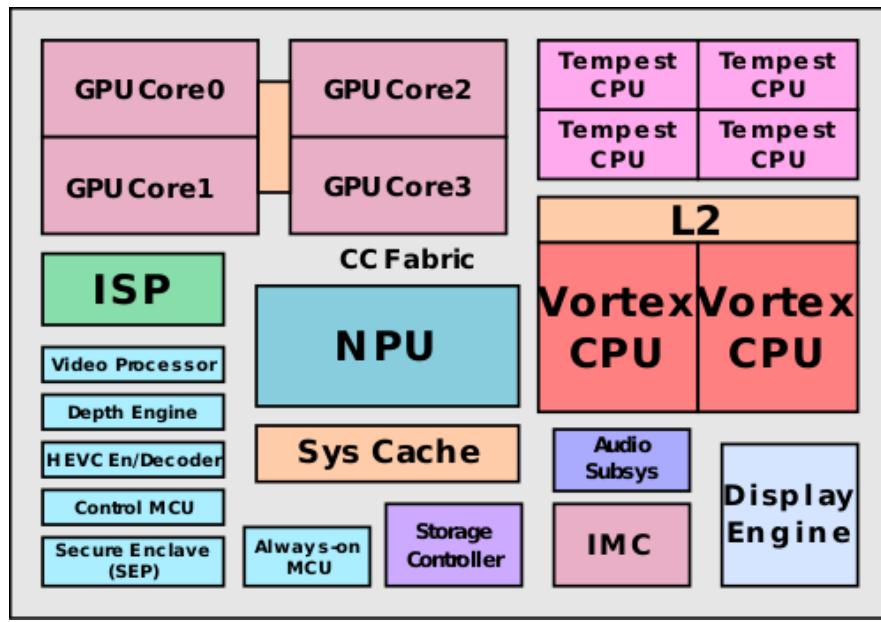
Apple M1 Pro



Special-purpose accelerators are increasingly used to improve performance & energy efficiency both in datacenters and at the edge

# Hardware Specialization in Mobile Chips

## System on chip (SoC)

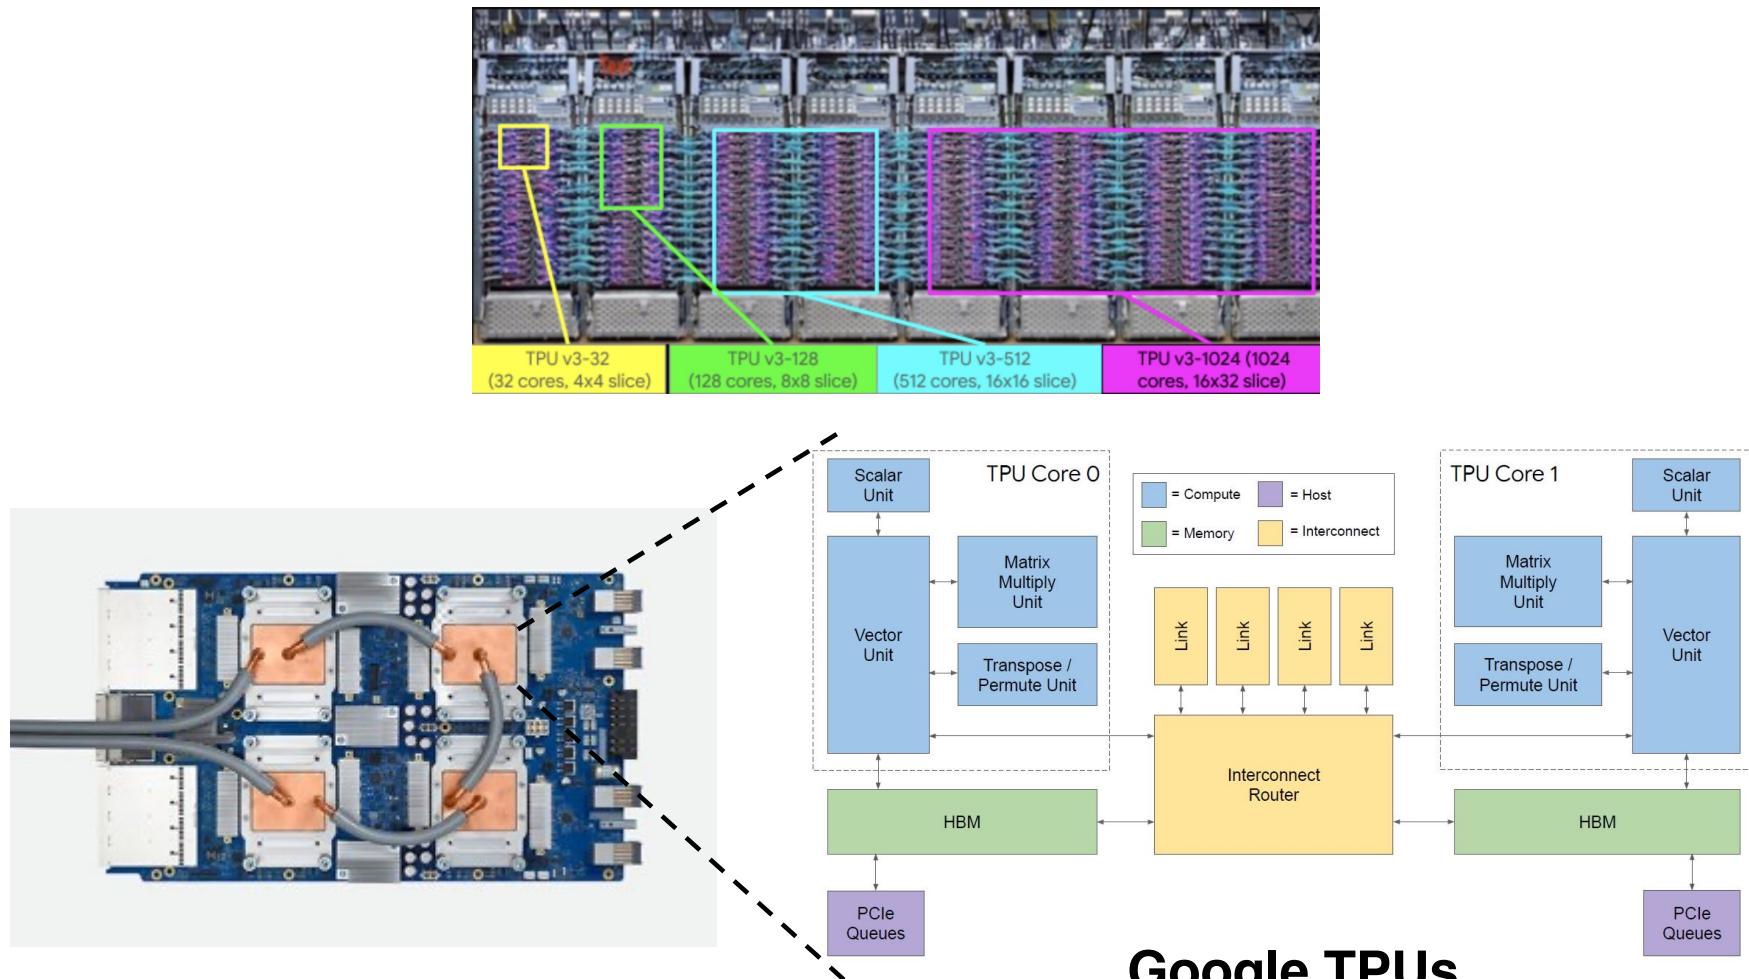


Apple 12 (iPhone X)

- ▶ Modern SoCs integrate a rich set of **special-purpose accelerators**
  - Speed up critical tasks
  - Reduce power consumption and cost
  - Increase energy efficiency

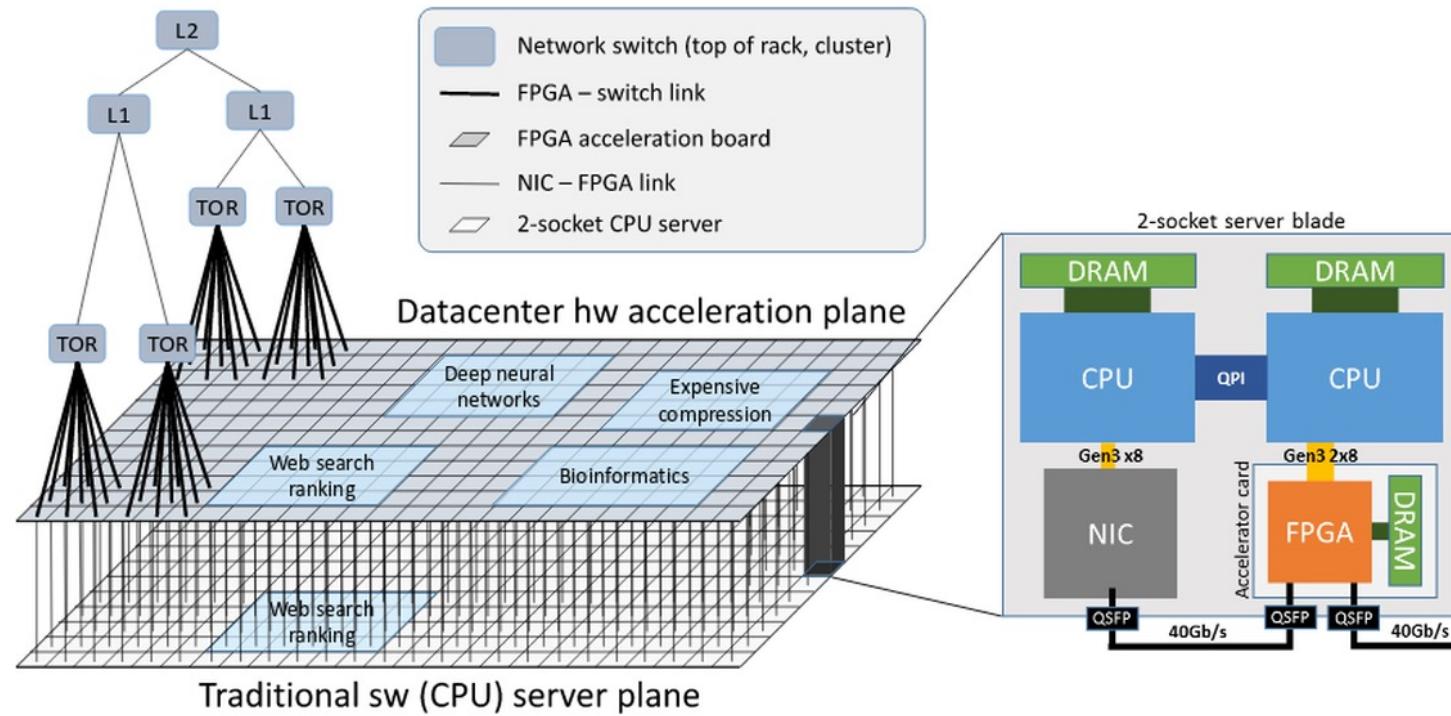
# Hardware Specialization in Datacenters

Application- and domain-specific accelerators are being deployed for a rich mix of compute-intensive applications in cloud datacenters



# Hardware Specialization in Datacenters

Application- and domain-specific accelerators are being deployed for a rich mix of compute-intensive applications in cloud datacenters

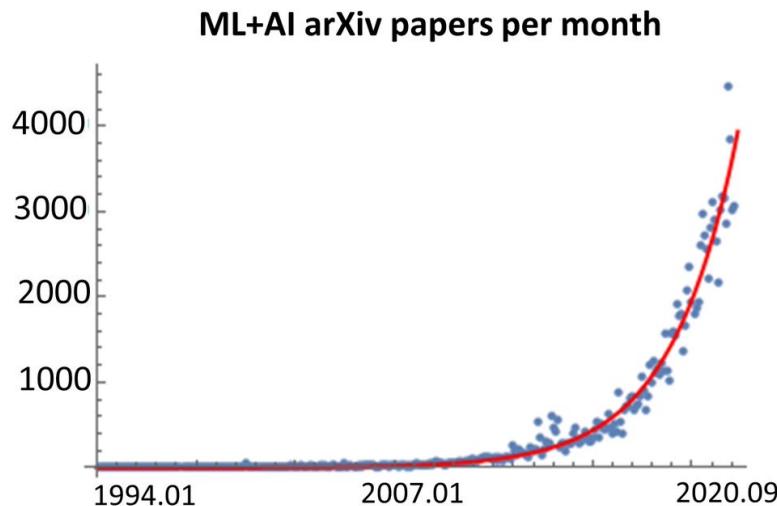


Microsoft Cloud FPGA Platforms

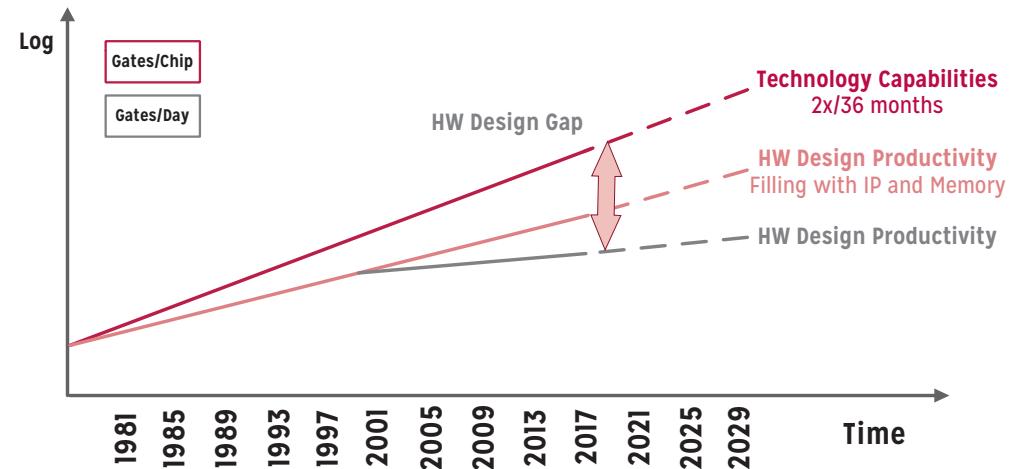
# Increasing Specialization Demands (Even) Higher Design Productivity

Can custom hardware evolve fast enough to keep up?

- ▶ Target of specialization is moving rapidly



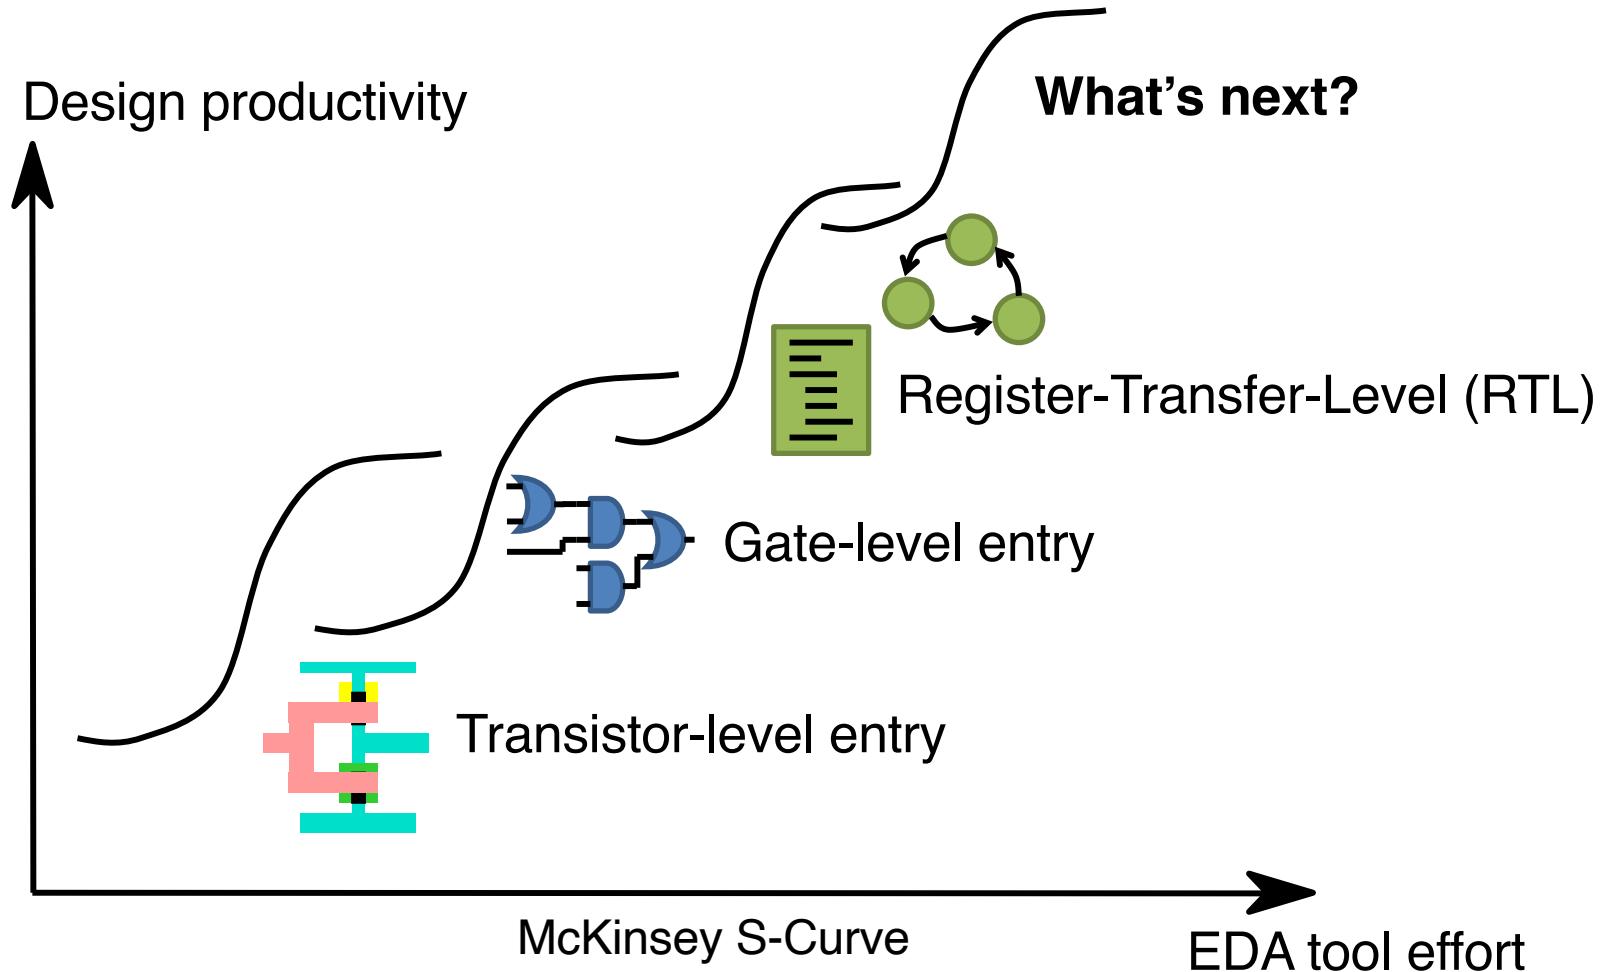
**Number of machine learning papers published on arXiv has outpaced Moore's Law**  
[Dean et al., IEEE Micro 2018]



**The design productivity gap**

[Source: Workshops on Extreme Scale Design Automation: Challenges and Opportunities for 2025 and Beyond]

# Evolution of Design Abstraction



[source: Kurt Keutzer, UCB]

# Motivation for High-Level Synthesis (HLS)

```
module dut(rst, clk, q);
    input rst;
    input clk;
    output q;
    reg [7:0] c;

    always @ (posedge clk)
    begin
        if (rst == 1b'1) begin
            c <= 8'b00000000;
        end
        else begin
            c <= c + 1;
        end
    end

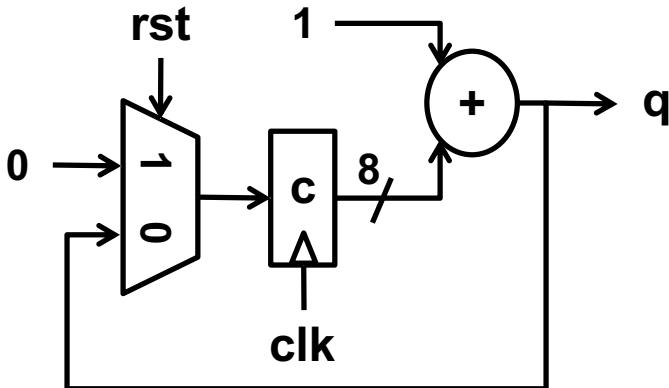
    assign q = c;
endmodule
```

RTL Verilog

VS.

```
uint8 dut() {
    static uint8 c;
    c+=1;
}
```

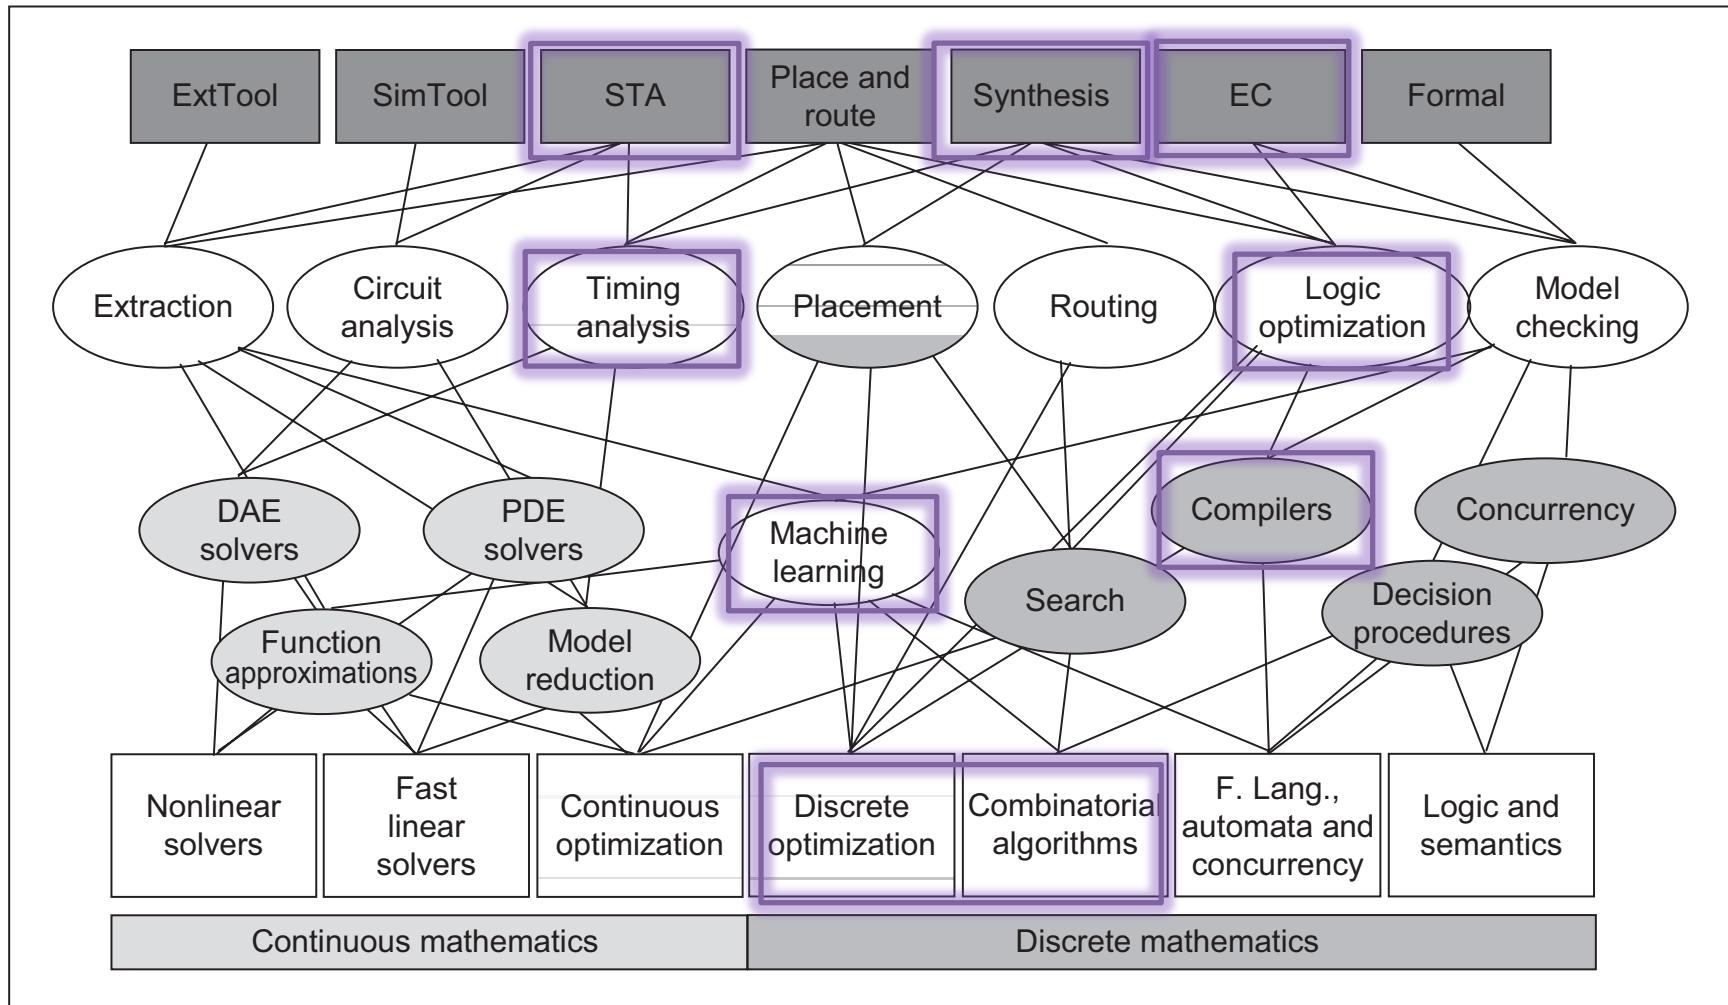
Automated  
with HLS



An 8-bit counter

# Algorithms Drive Automation

Topics touched on in 6775



## Key Algorithms in EDA

[source: Andreas Kuehlmann, Synopsys Inc.]

# Course Organization

- ▶ Refer to [syllabus](#) for course organization details

**Course Syllabus**  
**ECE 6775 High-Level Digital Design Automation**  
Fall 2024, Tuesday and Thursday 11:40am-12:55pm, Phillips 407

---

**1. Course Information**

**Lectures:** TuTh 11:40am-12:55pm, 407 Phillips Hall  
**Website:** <http://www.csl.cornell.edu/courses/ece6775>  
**CMS:** <https://cmsx.cs.cornell.edu>  
**Ed:** <https://edstem.org/us/courses/62024>

**Instructor:** Zhiru Zhang, [zhiruz@cornell.edu](mailto:zhiruz@cornell.edu)  
**Office Hours:** Thursday 4:30-5:30pm, Online

**Course Texts:**  
o Lecture slides/notes on course website  
o R. Kastner, J. Matai, and S. Neuendorffer, Parallel Programming for FPGAs, arXiv, 2018.  
o G. De Micheli, Synthesis and Optimization of Digital Circuits, McGraw-Hill, 1994.

**Supplementary Materials:**  
o S. Dasgupta, C.H. Papadimitriou, and U.V. Vazirani, Algorithms, McGraw-Hill, 2007.  
[link to online draft]  
o Additional reference papers will be posted as a course reader.

**2. Course Description and Objectives**

Targeted specialization of functionality in hardware has become arguably the best means to achieving improved compute performance and energy efficiency for a plethora of emerging applications. Unfortunately, it is a very unproductive practice to design and implement special-purpose accelerators using the conventional register transfer level (RTL) methodology. For this reason, both academia and industry are seeing an increasing use of high-level synthesis (HLS) to automatically generate hardware accelerators from software programs.

# Course Roadmap

- ▶ Lectures and paper\* discussions
  - **Background**
    - Introduction
    - Hardware specialization
    - Algorithm basics
    - FPGA\*
  - **High-level synthesis (HLS)**
    - C-based HLS
    - Front-end compilation
    - Scheduling
    - Resource sharing
    - Pipelining
  - **Advanced topics**
    - Deep learning acceleration\*
    - Domain-specific programming\*

# Preferred Background

- ▶ Working knowledge of the following at undergraduate level
  - C/C++
  - Digital logic and basic computer architecture concepts (e.g., adders, clock, registers, pipelining)
- ▶ Experiences with the following would increase appreciation & productivity
  - Algorithms and data structures
  - RTL design for FPGA or ASIC

## A Quick Poll



<http://pollev.com/ece6775>

**Sign in or register using your Cornell email ([NetID@cornell.edu](mailto:NetID@cornell.edu))**

Don't respond as a guest, or your answer won't be graded

# Learning Outcomes: The Tangibles

- ▶ High-level digital **design** methodologies
  - Design realistic accelerators *above RTL* using HLS design flow
  - Optimize accelerator performance using both manual source-level transformations and automatic HLS compiler optimizations
- ▶ High-level **design automation** algorithms
  - Fundamental compilation & synthesis techniques in HLS
    - e.g., SSA, scheduling, pipelining, resource sharing
  - Useful combinatorial optimization algorithms
    - e.g., graph algorithms, dynamic programming, greedy algorithms, integer linear programming

# Learning Outcomes: The Intangibles

- ▶ Develop a principled approach to analyzing accelerator design process and essential design factors (e.g., parallelism, resources, precision)
- ▶ Gain comprehensive insights into accelerator design from the perspective of an HLS compiler

We aim to achieve these objectives through **a balanced mix of theoretical foundations** (lectures & homework) and **practical applications** (labs & project)

# NOT Our Goals

- ▶ Teach you the design of microprocessors
- ▶ Cover the whole breadth of EDA
- ▶ Write RTL code
  - but this course will still improve your hardware design skills
- ▶ Make you an *expert* FPGA programmer
  - we use FPGA as a platform to prototype the accelerators

# Assignments

- ▶ Two problem sets (8%)
- ▶ Four lab assignments (22%)
  - Design & programming assignments leveraging high-level synthesis tools and software compilers
  - Experiments to be conducted on **eclinux** servers
    - % ssh -X <netid>@eclinux-01.ece.cornell.edu
    - Necessary tools will be installed in common directories

# Quizzes and Paper Readings

- ▶ Quizzes (6%)
  - You will need to answer pop quiz questions in most lectures
  - TWO lowest scores will be dropped
- ▶ Paper Readings (5%)
  - Two reading sessions
  - You are expected to read the paper or book chapter before the lecture, answer (more) quiz questions, and participate in discussions
  - Reading assignment will be announced at least one week in advance

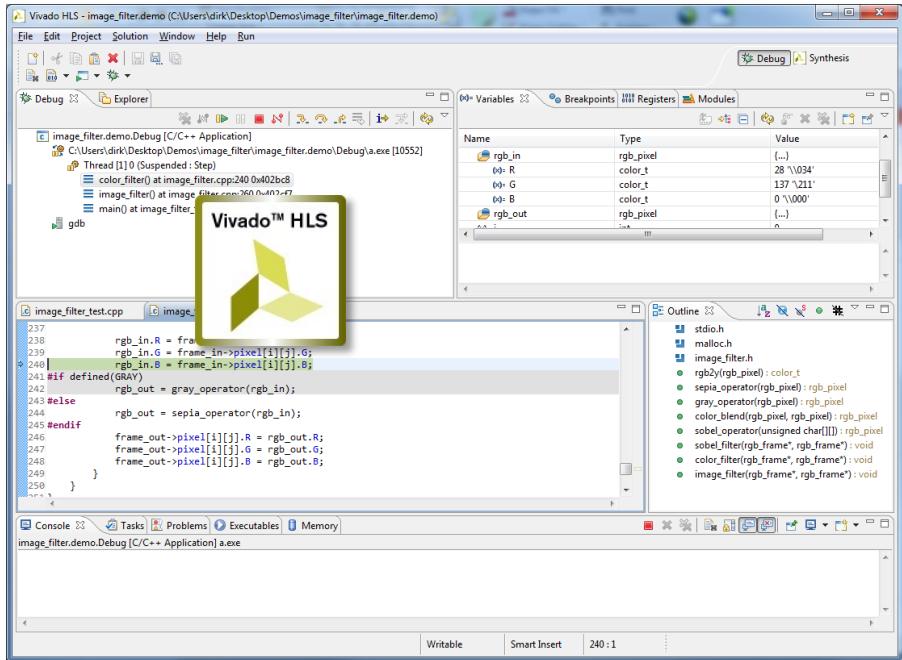
# Exam

- ▶ In-class midterm (20%)
  - Open notes & open book
  - When: Thursday October 17th
  - No sit-down final

## Final Project – 35%

- ▶ In-depth exploration of a research topic
  - (1) Designing new application-specific accelerators with HLS; OR
  - (2) Devising new automation algorithms/tools
  - 3-4 students / team, depending on class size
- ▶ Timeline
  - Proposal due after midterm
  - Weekly meeting with the instructor to track progress
  - Demo before the final week
  - Final report due by the final exam date

# High-Level Synthesis Tool



```
(vivado19) nz264@brg-zhang-xcel:~/shared/ece5997/mvmul-tutorial$ vivado_hls -f run.tcl
***** Vivado(TM) HLS - High-Level Synthesis from C, C++ and SystemC v2019.2.1 (64-bit)
**** SW Build 2729494 on Thu Dec 5 04:48:12 MST 2019
**** IP Build 2729494 on Thu Dec 5 07:38:25 MST 2019
** Copyright 1986-2019 Xilinx, Inc. All Rights Reserved.

source /opt/xilinx/Xilinx_Vivado_vitis_2019.2/Vivado/2019.2/scripts/vivado_hls/hls.tcl -notrace
INFO: Applying HLS Y2K22 patch v1.2 for IP revision
INFO: [HLS 200-10] Running '/opt/xilinx/Xilinx_Vivado_vitis_2019.2/Vivado/2019.2/bin/unwrapped/lnx64.o/vivado_hls'
INFO: [HLS 200-10] For user 'nz264' on host 'en-ec-brg-zhang-xcel.coecis.cornell.edu' (Linux_x86_64 version 3.10.0-1160.71.1.e
l7.x86_64) on Mon Aug 22 11:07:33 EDT 2022
INFO: [HLS 200-10] On OS 'Ubuntu 20.04 LTS' Linux release 7.9.2009 (Core)'
INFO: [HLS 200-10] In directory '/work/shared/users/phd/nz264/ece5997/mvmul-tutorial'
Sourcing Tcl script 'run.tcl'
INFO: [HLS 200-10] Opening and resetting project '/work/shared/users/phd/nz264/ece5997/mvmul-tutorial/mvmul_vitis.prj'.
INFO: [HLS 200-10] Adding design file 'mvmul_unroll.c' to the project
INFO: [HLS 200-10] Adding test bench file 'mvmul_top.c' to the project
INFO: [HLS 200-10] Opening and resetting solution '/work/shared/users/phd/nz264/ece5997/mvmul-tutorial/mvmul_vitis.prj/solutio
n1'.
INFO: [HLS 200-10] Cleaning up the solution database.
INFO: [HLS 200-10] Setting target device to 'xc7z020-clg484-1'
INFO: [SYN 201-201] Setting up clock 'default' with a period of 10ns.
INFO: [SCHED 204-61] Option 'relax_il_for_timing' is enabled, will increase II to preserve clock frequency constraints.
INFO: [HLS 200-10] Analyzing design file 'mvmul_unroll.c' ...
INFO: [HLS 200-111] Finished Linking Time (s): cpu = 00:00:11 ; elapsed = 00:00:18 . Memory (MB): peak = 1057.715 ; gain = 527
.219 ; free physical = 97063 ; free virtual = 219377
INFO: [HLS 200-111] Finished Checking Pragmas Time (s): cpu = 00:00:11 ; elapsed = 00:00:18 . Memory (MB): peak = 1057.715 ; g
ain = 527.219 ; free physical = 97063 ; free virtual = 219377
INFO: [HLS 200-111] Starting Standard Transformations ...
INFO: [HLS 200-111] Finished Standard Transformations Time (s): cpu = 00:00:12 ; elapsed = 00:00:19 . Memory (MB): peak = 1057.715
; gain = 527.219 ; free physical = 97048 ; free virtual = 219361
INFO: [HLS 200-10] Checking synthesizability ...
INFO: [HLS 200-111] Finished Checking Synthesizability Time (s): cpu = 00:00:12 ; elapsed = 00:00:19 . Memory (MB): peak = 105
7.715 ; gain = 527.219 ; free physical = 97063 ; free virtual = 219377
INFO: [HLS 200-489] Unrolling loop 'ACC_LOOP' ('mvmul_unroll.c:17') in function 'mvmul' completely with a factor of 16.
INFO: [XFORM 203-11] Balancing expressions in function 'mvmul' ('mvmul_unroll.c:6')...15 expression(s) balanced.
INFO: [HLS 200-111] Finished Pre-synthesis Time (s): cpu = 00:00:12 ; elapsed = 00:00:19 . Memory (MB): peak = 1057.715 ; gain
= 527.219 ; free physical = 97044 ; free virtual = 219358
INFO: [HLS 200-111] Finished Architecture Synthesis Time (s): cpu = 00:00:12 ; elapsed = 00:00:19 . Memory (MB): peak = 1057.7
15 ; gain = 527.219 ; free physical = 97043 ; free virtual = 219357
INFO: [HLS 200-10] Starting hardware synthesis ...
INFO: [HLS 200-10] Synthesizing 'mvmul' ...
WARNING: [SYN 201-107] Renaming port name 'mvmul/output' to 'mvmul/output_r' to avoid the conflict with HDL keywords or other
object names.
INFO: [HLS 200-10] _____
INFO: [HLS 200-10] _____ Implementing module 'mvmul'
INFO: [HLS 200-10] _____
INFO: [SCHED 204-11] Starting scheduling ...
INFO: [SCHED 204-11] Finished scheduling.
```

Tutorial on AMD Xilinx **Vivado HLS** v2019.2, Tuesday 9/10  
(we will be using the command-line interface in this course)

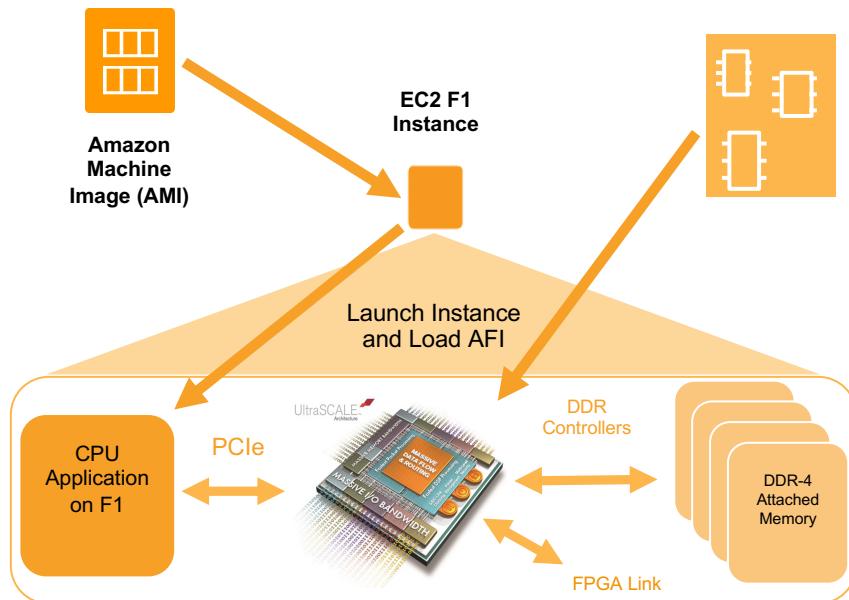
# Local Cluster of Embedded FPGAs

- ▶ For labs and project, we will use Zynq-based FPGA development boards (ZedBoard)
  - An FPGA SoC with a dual-core ARM CPU, which boots Linux



# Datacenter FPGA Platforms

- For the final project, students can also choose to explore datacenter FPGA platforms such as AMD Xilinx Alveo U280 and AWS F1 cloud instances



## Takeaway Points

- ▶ End of Dennard scaling leads to increasing **hardware specialization** to sustain improvement in performance and energy efficiency
- ▶ Increasing specialization and continued exponential growth in silicon capacity demands higher level of **design abstraction**
- ▶ HLS is a promising next step for EDA, which is fueled by sophisticated and yet scalable **algorithms**

# Before Next Lecture

- ▶ Action items
  - Check out the course website
  - **Read through the course syllabus**
  - **Verify your login on ecelinux**
    - ssh -X <netid>@ecelinux-01.ece.cornell.edu

# Acknowledgements

- ▶ These slides contain/adapt materials developed by
  - Prof. Jason Cong (UCLA)
  - Prof. David Z. Pan (UT Austin)